

Quantum error correction

September 2, 2018

1 Introduction

In the theoretical model of a universal quantum computer, quantum states are manipulated by quantum gates, i.e. unitary transformations, and eventually measurements are done to obtain the results of a quantum algorithm. In this model, the evolution of the quantum states over time is subject only to the quantum gates that make up this algorithm, and a quantum computer is usually described as a closed quantum system.

In reality, the situation is more challenging. The state of a quantum computer is constantly interacting with the environment, an interaction which we can try to minimize, but not fully suppress. This will imply changes to the state of the system that are beyond our control. In addition, as every measurement, also the measurements in a quantum computer will be subject to errors, and our realization of the quantum gates will not be perfect, but will be approximations to the ideal gate.

All these factors introduce errors into a quantum computer. For an individual gate, this error might still be small, but if want to build an universal quantum computer that executes algorithms consisting of thousands or millions of gates, the errors are going to propagate and will soon render the result unusable. So we will have to detect and control these errors.

This situation is not entirely new. Also in classical computers, bits in memory or on a hard drive can flip due to some unwanted interaction with the environment, and computer scientists have developed error correction methods to suppress these errors. Typically, these mechanisms create one or more copies of the state and, after some time, compare the copies. If a bit in one of the copies has flipped, chances are that the other bits remain intact, and we can detect the error by comparing the bits.

However, in the world of quantum mechanics, the situation is more complicated. In fact, it is not possible to create a copy of a quantum state, a fact that is known as the **No-cloning theorem** (see [8] and [9]).

Let us see - following the argument given in chapter 12 of [1] why this is true. Suppose we had a device that is able to create a copy of a quantum state $|\psi\rangle$. More formally, this device would act on a system consisting of two subsystems. We would initialize the first subsystem in the state $|\psi\rangle$ and the second subsystem in some initial state $|e\rangle$. Then the device would operate and leave us with a system in which the second subsystem is in the state $|\psi\rangle$ as well.

As any time evolution of a quantum system allowed by the laws of quantum mechanics, this device would be described by a unitary operator, acting on the

combined system. This operator, let us call it U , would therefore act on the states according to the rule

$$U(|\psi\rangle|e\rangle) = |\psi\rangle|\psi\rangle$$

where we use the usual notation for a tensor product. This should be true for any state $|\psi\rangle$. Therefore, given any two states $|\psi\rangle$ and $|\Phi\rangle$, we would have the relations

$$\begin{aligned} U(|\psi\rangle|e\rangle) &= |\psi\rangle|\psi\rangle \\ U(|\Phi\rangle|e\rangle) &= |\Phi\rangle|\Phi\rangle \end{aligned}$$

Let us now compute the scalar product of the two states on the right hand side. First, we can do this directly and obtain

$$\langle|\psi\rangle|\psi\rangle, |\Phi\rangle|\Phi\rangle\rangle = \langle\psi, \Phi\rangle\langle\psi, \Phi\rangle$$

On the other hand, we could use the fact that the operator U is unitary to equate the tensor product with the tensor product of the states on the left hand side of our equation, and obtain

$$\langle|\psi\rangle|\psi\rangle, |\Phi\rangle|\psi\rangle\rangle = \langle|\psi\rangle|e\rangle, |\Phi\rangle|e\rangle\rangle = \langle\psi, \Phi\rangle\langle e, e\rangle$$

Without loss of generality, we can assume that the state $|e\rangle$ is normalized. Then comparing these two expressions yields the equality

$$\langle\psi, \Phi\rangle = \langle\psi, \Phi\rangle\langle\psi, \Phi\rangle$$

This is only possible if either the two states $|\psi\rangle$ and $|\Phi\rangle$ are orthogonal or if $\langle\psi, \Phi\rangle = 1$. But of course we can easily find examples of states for which neither of these relations hold and the equality is violated. Thus our device will never be able to clone an arbitrary state.

There is a second reason why simple error correction mechanisms as they are known from the world of classical computing will in general fail in the quantum world - these mechanisms perform measurements. However, in a quantum computer, a measurement will typically destroy our superposition and therefore affect the outcome of a quantum algorithm. So we need a smarter approach.

2 The three bit code

Let us now investigate a simple example that demonstrates that it is still possible to account for errors in a quantum setup. Suppose that we are looking at a communication channel over which we are somehow able to transmit individual qubits. Suppose further that the only error that can occur during this transmission is a bit flip, i.e. that with a probability p , the states $|0\rangle$ and $|1\rangle$ are switched during the transmission (this is of course not realistic, and we will investigate other sources of errors below). Using the standard notation

$$X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$




Figure 1: Three qubit encoding

for the bit-flip operator, a state $|\psi\rangle$ will therefore be converted into the state $X|\psi\rangle$ with probability p .

Now suppose that we encode every single qubit state in a three-qubit state before transmitting it. Thus we use the following encoding

$$\begin{aligned} |0\rangle &\mapsto |000\rangle \\ |1\rangle &\mapsto |111\rangle \end{aligned}$$

It is not difficult to see that this encoding can in fact be realized by a unitary circuit - the circuit in figure 1 will do the trick.

Let us see why this works. If $|\psi\rangle = |0\rangle$, the two CNOT gates do nothing, so the outcome is $|000\rangle$. If $|\psi\rangle = |1\rangle$, both CNOT gates will flip the respective qubit, so the outcome is $|111\rangle$. Thus a general state $|\psi\rangle = a|0\rangle + b|1\rangle$ will be encoded as

$$a|000\rangle + b|111\rangle$$

Note that this is not equal to the tensor product $|\psi\rangle|\psi\rangle|\psi\rangle$ and therefore this circuit does not violate the no-cloning theorem.

We now send each of the three qubits through our channel independently. We then perform an operation called a *syndrome measurement*. Specifically, we apply the circuit displayed in figure 2.




Figure 2: Syndrome measurement

This circuit uses two ancilla qubits to effectively compare the parity of the individual qubits transmitted by the channel. To see this, note that the combination of the first two CNOT gates act on the same ancilla qubit. If the first two transmitted qubits are both $|0\rangle$, the ancilla is unchanged. But this is also true if both qubits are $|1\rangle$, as then the ancilla is inverted twice. So the first ancilla qubit will compare the first and second transmitted qubit, and the second ancilla qubit compares the second and third qubit.

Now suppose that at most one qubit has been flipped during the transmission, i.e. during the transmission, one of the following four unitary operators has been applied to our initial state $|\psi\rangle = a|000\rangle + b|111\rangle$.

$$\begin{aligned} &1 \otimes 1 \otimes 1 \\ &X \otimes 1 \otimes 1 \\ &1 \otimes X \otimes 1 \\ &1 \otimes 1 \otimes X \end{aligned}$$

Let us now write down the state of the overall system after performing the syndrome measurement, i.e. after adding the two ancilla qubits and applying the syndrome measurement circuit. We obtain

$$\begin{aligned} &(a|000\rangle + b|111\rangle)|00\rangle \\ &(a|100\rangle + b|011\rangle)|11\rangle \\ &(a|010\rangle + b|101\rangle)|10\rangle \\ &(a|001\rangle + b|110\rangle)|01\rangle \end{aligned}$$

We can now measure the two ancilla bits. This measurement will not change the state of the three transmitted qubits, but tell us which of the four cases we are in! Also note that if we measure the ancilla qubits, we will obtain the result with probability one in each of the four cases. Thus after performing

the measurement, we know with certainty which of the four operators has been applied during the transmission and can perform a correction by flipping the respective bit once more. Of course this breaks down if more than one qubit is flipped, but as the probability for this to happen is smaller than p , we have been able to reduce the error rate.

So what we have learned so far is that it is possible, also in the quantum world, to detect and protect against pure bit flip errors without destroying the superposition. But there is more we can learn from that example. In fact, let us revisit our original assumption that the only thing that can go wrong is that the operator X is applied to some of the qubits and allow a more general operator. Suppose, for instance, that our error is represented by applying to at most one of our qubits the operator

$$E = (1 - \epsilon)1 + \epsilon X$$

for a small ϵ . This operator is of course in general not unitary, and thus the operator represents the impact of decoherence, i.e. an interaction with the environment (see section 5.1 of [10] for a discussion of the most general type of errors that we have to deal with). Suppose we apply this operator to the first qubit during the transmission. Then our state after the transmission will be

$$(1 - \epsilon)(\alpha|000\rangle + \beta|111\rangle) + \epsilon(\alpha|100\rangle + \beta|011\rangle)$$

After applying the error syndrome measurement circuit, this state will again be in a superposition, namely

$$(1 - \epsilon)(\alpha|000\rangle + \beta|111\rangle)|00\rangle + \epsilon(\alpha|100\rangle + \beta|011\rangle)|11\rangle$$

Assuming that our original state was normalized, the norm of this state will be $(1 - \epsilon)^2 + \epsilon^2 = 1 - 2\epsilon$. Now we measure the syndrome. As always in quantum mechanics, measurement corresponds to projection onto an eigenspace. Thus, after the measurement, the state will - up to normalization - be one of the states

$$\begin{aligned} &(\alpha|000\rangle + \beta|111\rangle)|00\rangle \\ &(\alpha|100\rangle + \beta|011\rangle)|11\rangle \end{aligned}$$

and the probabilities to be in these states will be $\frac{1-\epsilon}{\sqrt{1-2\epsilon}}$ and $\frac{\epsilon}{\sqrt{1-2\epsilon}}$. Thus the measurement will collapse the state into one of two states, where one state corresponds to no error at all, and one state corresponds to a bit flip error. In other words, whereas the actual error that was applied was a linear combination of the identity and a bit flip error, the process of detecting and correcting the error somehow forces the system to make a decision: either there is a full bit-flip error, or there is no error at all. This fact is called the *digitization of errors* (see also [10] for a more complete treatment). Thus instead having to prepare for a continuum of error operators, it suffices to fix a discrete set of errors - which obviously is of utmost importance for practical implementations.

3 Complete error codes

So far we have worked with a code which is able to protect against a bit flip error. But of course this is not the only type of error that can occur. At the

first glance, it looks like there is a vast universe of potential errors that we have to account for, as in theory, the error could be any unitary operator. However, using arguments similar to the discussion in the last section, one can show that it suffices to protect against two types of errors: the bit flip error discussed above and the *phase flip error*, represented by the matrix

$$Z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$$

The first error correction code that was able to handle bit flip errors and phase flip errors (and combinations thereof) and therefore any possible type of error was invented in 1995 by P. Shor in [12]. To understand this code, it is useful to first design a code that protects against phase flip errors alone. The idea for this code is based on the relations

$$\begin{aligned} Z|+\rangle &= |-\rangle \\ Z|-\rangle &= |+\rangle \end{aligned}$$

where we denote by $|+\rangle$ and $|-\rangle$ the alternative basis of a one-qubit Hilbert space given by

$$\begin{aligned} |+\rangle &= \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) = H|0\rangle \\ |-\rangle &= \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle) = H|1\rangle \end{aligned}$$

Here H is the Hadamard operator which maps the standard basis to the basis given by these two vectors and is therefore described by the unitary matrix

$$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$$

In other words, applying a phase flip error is the same as applying the Hadamard operator, followed by a bit flip followed again by the Hadamard operator, i.e.

$$HXH = Z$$

This implies that a phase flip and a bit flip are two different views of the same thing and therefore that a code that is able to detect and correct bit flip errors is also able to correct phase flip errors if we apply the Hadamard matrix to convert back between these two views. Specifically, we can encode a state by applying the three bit code followed by the Hadamard operator on each qubit. Thus

$$\begin{aligned} |0\rangle &\mapsto |+++ \rangle \\ |1\rangle &\mapsto |-- - \rangle \end{aligned}$$

For the error syndrome measurement, we apply the Hadamard operator on each of the three qubits, followed by the circuit shown in figure 2, followed again by the Hadamard operator applied to each qubit. The values of the two ancilla qubits then again tell us whether a phase flip has occurred and for which qubit.

If, for instance, a phase flip is applied to the first qubit, a state $a|0\rangle + b|1\rangle$ becomes

$$a| - ++ \rangle + b| + -- \rangle$$

Applying the Hadamard operator $H^{\otimes 3}$ then turns this state into the state that would be the result of a bit flip error, i.e. into

$$a|100\rangle + b|011\rangle$$

When we now apply our parity measurement circuit in figure 2, we obtain the state

$$(a|100\rangle + b|011\rangle)|11\rangle$$

Finally, we apply the Hadamard operator again to each of the three primary qubits, i.e. we apply $H^{\otimes 3}$, and obtain

$$(a| - ++ \rangle + b| + -- \rangle)|11\rangle$$

We can now again measure the two ancilla qubits and use the result to distinguish between the four different error situations that can arise if we assume that at most one phase flip can take place, and use this information to apply the necessary corrections.

With this understanding, we are now ready to describe Shor's code that corrects phase flip and bit flip errors (and their combination ZX). The idea of that algorithm is to concatenate the encoding that protects against phase flip errors with the encoding that protects against bit flip errors.

Specifically, we encode our states as follows. We first encode once using the phase flip code. Thus, we map $|0\rangle$ to $| + ++ \rangle$ and $|1\rangle$ to $| - -- \rangle$. We then apply the bit flip code to each of the resulting qubits. The bit flip code maps $|0\rangle$ to $|000\rangle$ and $|1\rangle$ to $|111\rangle$, and therefore maps the linear combination $|+\rangle$ to $\frac{1}{\sqrt{2}}(|000\rangle + |111\rangle)$ and similarly $|-\rangle$ to $\frac{1}{\sqrt{2}}(|000\rangle - |111\rangle)$. Thus the full encoding uses nine bits and is given by

$$\begin{aligned} |0\rangle &\mapsto \frac{1}{2\sqrt{2}}(|000\rangle + |111\rangle)(|000\rangle + |111\rangle)(|000\rangle + |111\rangle) \\ |1\rangle &\mapsto \frac{1}{2\sqrt{2}}(|000\rangle - |111\rangle)(|000\rangle - |111\rangle)(|000\rangle - |111\rangle) \end{aligned}$$

Now suppose a bit flip error occurs on this encoded state, say on the first qubit. The second and third block of three qubits each are then unaffected, and the error to the first block corresponds to an error that occurs after applying a bit flip error correction code to the three bits in this block. It can therefore be detected and corrected as before.

Now let us see what happens if a phase flip occurs. First, let us introduce a shorthand notation. We abbreviate the three qubit states that appear in the nine bit encoding as

$$\begin{aligned}|A\rangle &= |000\rangle + |111\rangle \\ |B\rangle &= |000\rangle - |111\rangle\end{aligned}$$

With this notation, we can now simplify the expressions for the nine qubit encoding as follows.

$$\begin{aligned}|0\rangle &\mapsto \frac{1}{2\sqrt{2}}|AAA\rangle \\ |1\rangle &\mapsto \frac{1}{2\sqrt{2}}|BBB\rangle\end{aligned}$$

Now let us calculate the result of applying the Hadamard operator to each block of the encoded state. A short calculation shows that

$$\begin{aligned}H|A\rangle &= H(|000\rangle + |111\rangle) = \frac{1}{\sqrt{2}}(|000\rangle + |011\rangle + |110\rangle + |101\rangle) \\ H|B\rangle &= H(|000\rangle - |111\rangle) = \frac{1}{\sqrt{2}}(|111\rangle + |001\rangle + |100\rangle + |010\rangle)\end{aligned}$$

In the first line, each state contains an even numbers of ones, in the second line each state contains an odd number of ones. Thus if we take an ancillary qubit with initial value $|0\rangle$ and apply a sequence of three CNOT gates to it, controlled by the three qubits, the result will be $|0\rangle$ for the first line and $|1\rangle$ for the second line. Thus this circuit will allow us to tell the state $|A\rangle$ and $|B\rangle$ apart.

Now let us take a look at the circuit in figure 3 that combines six of these detection circuits.

Repeating the arguments above for each of the blocks, we see that if we apply this circuit to the states $|AAA\rangle$ and $|BBB\rangle$, both ancillary qubits will be zero. Now suppose that a phase flip in one of the first qubits occurs in, say, the state $|AAA\rangle$. Clearly, this phase flip will result in the state $|BAA\rangle$. When we run the circuit on this state, the first three CNOT gates applied to the first ancillary qubit will give $|1\rangle$. The second set of three CNOT gates will not modify this state any more, so the first ancillary qubit will be $|1\rangle$. The same holds for the second ancillary qubit, and so the two ancillary qubits will be $|11\rangle$. If a phase flip occurs in the second block, we get the state $|ABA\rangle$. Now the first three CNOT gates leave the first ancillary qubit in state $|0\rangle$, whereas the second three CNOT gates invert it, so that the first ancillary qubit will be $|1\rangle$. However, the second ancillary qubit will remain in state $|0\rangle$, so that the combined state will be $|10\rangle$. Finally, the state $|AAB\rangle$ which is the result of phase flip in the third block will give us $|01\rangle$. A similar reasoning applies to a phase flip on the state $|BBB\rangle$. The upshot of this discussion is that very similar to the three qubit code, the value of the two ancillary qubits can be measured and will tell us in which block a phase flip occurred. Changing the phase of any of the qubits in this block will then correct the error. This shows that the Shor code can also be used to detect and correct phase flip errors, as long as only one bit flips.




Figure 3: Nine qubit code - syndrome measurement

Similarly, one can show that also a combination of phase flip and bit flip errors, i.e. the result of applying XZ , can be detected. Thus this code is able to detect all discrete errors and therefore - by the principle of the digitalization of errors - all errors as long as only one qubit is affected.

Let us now summarize some of the formal properties that this code has. First, the encoding can be described as carving out a 2-dimensional subspace C , called the *code space*, in the larger space spanned by all (in this case nine) qubits. More generally, the space C can have dimension 2^n and the full Hilbert space can have dimension 2^k , in which case we can encode n qubits in a larger set of k qubits (in the nine qubit code example, $k = 9$ and $n = 1$).

Hence the code space is spanned by a set of 2^n basis vectors $|\psi_i\rangle$ that we call *code words*. In the case $n = 1$, it is common practice to denote the codewords by $|0\rangle_L$ and $|1\rangle_L$ to indicate that they represent a *logical qubit* encoded using k *physical qubits*.

In addition, there is a set of *error operators*, i.e. a finite set of operators $\{E_\alpha\}$ acting on the larger Hilbert space. These operators represent the impact of discretized noise.

Now what are the conditions on C and the E_α that allow us to perform error correction? A first condition is that no matter which errors occur, we can still tell the code words apart. This is guaranteed if the various error operators map different code words onto mutually orthogonal subspaces, in other words if

$$\langle\psi_i|E_\alpha^\dagger E_\beta|\psi_j\rangle = 0$$

whenever $i \neq j$ and for all α, β . Thus, even in the presence of errors, the

different code words will never overlap.

What about the case that both code words are the same? It is tempting to ask for the condition

$$\langle \psi_i | E_\alpha^+ E_\beta | \psi_i \rangle = \delta_{\alpha\beta}$$

i.e. to require that different errors map the same codeword to orthogonal subspaces. This would make recovery very easy. We could perform the measurements that correspond to projections onto these subspaces to detect the error and correct them by applying the inverse of the error operator. However, this condition is too restrictive. In the case of the nine qubit code, for example, a phase flip Z_1 and a phase flip Z_2 in the second qubit both map the state $|AAA\rangle$ to the same state, namely $|BAA\rangle$. However, the same applies for the correction, i.e. we do not have to distinguish between these two errors as they act similarly on the code space. In [14], a more general condition is presented which captures this case (see also [3]):

$$\langle \psi_i | E_\alpha^+ E_\beta | \psi_j \rangle = C_{\alpha\beta} \delta_{ij} \quad (1)$$

for a hermitian matrix $C_{\alpha\beta}$. If the matrix has full rank, the code is called *non-degenerate*, otherwise - as in the case of the nine qubit code - the code is called *degenerate*.

For the Shor code, we have to use nine physical qubits to encode one logical qubit. In 1996, Steane discovered a code ([11]) that requires only seven physical qubits. In the same year, a code that requires only five qubits was presented ([15]) and it was shown that this is a lower bound, i.e. there is no error correction code that requires less than five qubits.

4 The stabilizer formalism

The formalism of stabilizer codes has been introduced by D. Gottesman in his thesis ([5]) and is very useful to identify and describe a large class of quantum error correction codes. To get the idea, let us first consider the action of products of Pauli matrices on n -qubit systems. In what follows, we will use an upper index to indicate the qubit on which an operator acts in a multi-qubit system. First, let us calculate

$$\begin{aligned} \sigma_z^1 \sigma_z^2 |00\rangle &= |00\rangle \\ \sigma_z^1 \sigma_z^2 |11\rangle &= |11\rangle \\ \sigma_z^1 \sigma_z^2 |10\rangle &= -|10\rangle \\ \sigma_z^1 \sigma_z^2 |01\rangle &= -|01\rangle \end{aligned}$$

Thus the hermitian operator $\sigma_z^1 \sigma_z^2$ is able to compare the signs of the two qubits. A state in which these two qubits have the same sign is in the eigenspace with eigenvalue one, a state in which the signs differ is in the eigenspace with eigenvalue minus one.

Looking again at the Shor nine qubit code, this implies that the process of detecting a bit flip error can be implemented using measurements corresponding

to these operators. In order to detect a bit flip error in the first block of three qubits, we can measure $\sigma_z^1 \sigma_z^2$ and $\sigma_z^1 \sigma_z^3$. The outcomes of these two measurements then correspond to the values of the ancillary qubits used in our syndrome measurement circuits and tell us whether a bit flip occurred and where - the same is true for the other blocks.

A similar observation holds for phase flip errors. In fact, it is obvious that - using the same notation as before - we have

$$\begin{aligned}\sigma_x^1 \sigma_x^2 \sigma_x^3 |A\rangle &= |A\rangle \\ \sigma_x^1 \sigma_x^2 \sigma_x^3 |B\rangle &= -|B\rangle\end{aligned}$$

Thus a measurement corresponding to the operator $\sigma_x^1 \sigma_x^2 \sigma_x^3$ can be used to distinguish between the states $|A\rangle$ and $|B\rangle$. Consequently, the syndrome measurement circuit for the phase flip error corresponds to measurements of the operators $\sigma_x^1 \sigma_x^2 \sigma_x^3 \sigma_x^4 \sigma_x^5 \sigma_x^6$ and $\sigma_x^1 \sigma_x^2 \sigma_x^3 \sigma_x^7 \sigma_x^8 \sigma_x^9$.

As described in [5], there is a total of eight operators M_1, \dots, M_8 of this type that we have to measure to be able to detect any possible one qubit phase flip or bit flip error. Each of these operator stabilizes the two logical one-qubit states $|0\rangle_L$ and $|1\rangle_L$ of the Shor code, i.e. these states are eigenstates with eigenvalue one. Also note that the matrices M_1, \dots, M_8 all commute and therefore span an abelian group.

Let us now try to formalize and generalize this observation. We consider a system with k qubits and the Pauli group $\mathcal{G} = \mathcal{G}_k$. Within the Pauli group, we consider an abelian subgroup S called the *stabilizer*. We can then look at the linear subspace T_S which is the intersection of all $+1$ eigenspaces of the elements in S , i.e.

$$T_S = \{|\psi\rangle \text{ such that } s|\psi\rangle = |\psi\rangle \forall s \in S\}$$

To avoid confusion, let us stress that we do not assume that S is the full stabilizer, i.e. there might be elements of \mathcal{G} that also keep all elements of T fixed but are not in S .

The idea of stabilizer codes is now to use a subspace of the form T_S as code space for a quantum error correction code. In general, the stabilizer S is described in terms of l independent generators, i.e. $S = \langle g_1, \dots, g_l \rangle$. One can then show (see A.2) that the dimension of the subspace T_S is equal to 2^{k-l} . In the example of the nine qubit code, we are considering the subgroup S generated by the elements M_1, \dots, M_8 which turn out to be independent. The overall Hilbert space has dimension $2^k = 512$ as $k = 9$. Therefore the subspace T_S has dimension $2^{9-8} = 2$ and therefore describes one logical qubit.

Now let us turn to the question which errors a code provided in this way by a stabilizer. Assume that we are given a set of error operators $\{E_\alpha\} \subset \mathcal{G}$ as before. For any two indices, we can then consider the operator $E_{\alpha\beta} = E_\alpha^\dagger E_\beta$. Given any $E_{\alpha\beta}$ and any operator $M \in S$, these two operators either commute or anti-commute, as they are in \mathcal{G} .

Let us first suppose that M and $E_{\alpha\beta}$ anti-commute. Then, given any two code words $|\psi_i\rangle$ and $|\psi_j\rangle$, we can calculate

$$\langle \psi_i | E_\alpha^\dagger E_\beta | \psi_j \rangle = \langle \psi_i | M E_{\alpha\beta} | \psi_j \rangle = -\langle \psi_i | E_{\alpha\beta} M | \psi_j \rangle = \langle \psi_i | E_\alpha^\dagger E_\beta | \psi_j \rangle$$

where we have used that $M|\psi_i\rangle = M^+|\psi_i\rangle = |\psi_i\rangle$ as M is either plus or minus M^+ . This is only possible if

$$\langle\psi_i|E_\alpha^+E_\beta|\psi_j\rangle = 0$$

This is in line with our earlier condition (1)! Similarly, if $E_{\alpha\beta}$ itself is in S , we obtain that

$$\langle\psi_i|E_\alpha^+E_\beta|\psi_j\rangle = \langle\psi_i|E_{\alpha\beta}|\psi_j\rangle = \langle\psi_i|\psi_j\rangle\delta_{ij}$$

which again matches condition (1). Also note that due to the fact that E_α is either hermitian or anti-hermitian, we have $E_\alpha^+E_\beta = \pm E_\alpha E_\beta$. Thus we have shown

THEOREM 4.1 (see [5]). *Suppose that S is an abelian subgroup of the Pauli group \mathcal{G} with a non-trivial invariant subspace T_S . If $\{E_\alpha\}$ is a set of error operators in \mathcal{G} such that for any two indices α, β , the operator $E_\alpha E_\beta$ either anti-commutes with some element of S or is itself in S , then T_S defines an error correction code that is able to correct the errors $\{E_\alpha\}$.*

Let us again consider the nine qubit Shor code as an example, and let $E_1 = \sigma_x^1$ and $E_2 = 1$. Then $E_{12} = \sigma_x^1$ which clearly anticommutes with $M_1 = \sigma_z^1\sigma_z^2$. Thus we can detect and correct single bit flip errors in the first qubit (and similarly in all other qubits). Similarly, if $E_1 = \sigma_z^1$, the operator E_{12} will anti-commute with the operator $M_7 = \sigma_x^1\sigma_x^2\sigma_x^3\sigma_x^4\sigma_x^5\sigma_x^6$. Thus we are also able to correct phase flip errors in the first qubit. A similar reasoning applies to all other one qubit errors, thus we again find that the Shor code can correct them all (see [5]).

To find the possible set of error operators, we therefore need to take a look at the centralizer $C(S)$ which - as common in group theory - is defined to be the set of group elements in \mathcal{G} that commute with every element of S . Now, in our case, the centralizer is in fact equal to the normalizer $N(S)$, i.e. the set of all group elements that fix S under conjugation.

To see this, let us first assume that we are given some element A of the Pauli group such that $A \in N(S)$. Then, for any $s \in S$, we have

$$A^+sA \in S$$

by definition of the normalizer. But being elements of the Pauli group, s and A either commute or anti-commute. We claim that they commute, which would prove that $A \in S$. In fact, if A and s anti-commute, then

$$A^+sA = -A^+As = -s$$

so $-s$ would be in S . But $s \in S$, so that this would imply that $-1 \in S$, which we have excluded. This proves that $N(s) \subset C(S)$. Conversely, it is clear that $C(S) \subset N(S)$.

In light of this identity, we can now rephrase the condition expressed in theorem 4.1: the code given by the stabilizer S corrects the error operators $\{E_\alpha\}$ if for any two indices E_α, E_β , we have

$$E_\alpha E_\beta \in S \cup (\mathcal{G} - N(S))$$

For the elements in the normalizer $N(S)$, there is a different story to tell. First, if $n \in N(S)$, the matrix n maps T_S to itself. In fact, if $|\psi\rangle \in T_S$, we have

$$n|\psi\rangle = ns|\psi\rangle = sn|\psi\rangle$$

for every $s \in S$. This shows that $n|\psi\rangle$ is fixed by all $s \in S$ and therefore in T_S . Thus the elements of $N(S)/S$ define an automorphism of T_S . We can interpret these matrices as quantum gates that now act on the *logical qubits* in the code space.

Let us consider an example to see how this works in general. We consider the following four generators, acting on a quantum register with five qubits (and therefore providing a two-dimensional invariant subspace T_S , i.e. one logical qubit).

$$\begin{aligned} S_1 &= \sigma_x^1 \sigma_z^1 \sigma_z^2 \sigma_z^4 \sigma_x^5 \sigma_z^5 \\ S_2 &= \sigma_x^2 \sigma_z^3 \sigma_z^4 \sigma_x^5 \\ S_3 &= \sigma_z^1 \sigma_z^2 \sigma_x^3 \sigma_x^5 \\ S_4 &= \sigma_z^1 \sigma_z^3 \sigma_x^4 \sigma_z^4 \sigma_x^5 \sigma_z^5 \end{aligned}$$

It is a bit tedious, but not difficult, to verify that these generators commute. For instance, if we move S_2 past S_1 , we pick up one minus one from the second qubit and one from the last qubit. When we commute S_3 past S_3 , we pick up a sign from the second qubit and a sign from the third qubit which is again an even number and so forth.

Let us now in addition introduce two additional operators that we call \bar{X} and \bar{Z} which are defined as follows.

$$\begin{aligned} \bar{X} &= \sigma_z^1 \sigma_z^4 \sigma_x^5 \\ \bar{Z} &= \sigma_z^1 \sigma_z^2 \sigma_z^3 \sigma_z^4 \sigma_x^5 \end{aligned}$$

Again, a short calculation shows that each of these operators commutes with all the S_i . However, \bar{X} and \bar{Z} do not commute, but anti-commute! Thus the algebra of these two matrices is exactly as that of σ_x and σ_z . In other words, we obtain an action of the Pauli group \mathcal{G}_1 on the space T_S , i.e. on the logical qubit described by this code.

Looking at this code, it appears to be an extremely clever combination of Pauli matrices which is nearly impossible to find. However, the real beauty of the stabilizer formalism is that it reduces the problem of finding a code like this one - and the operators \bar{X} and \bar{Z} - to linear algebra over the field \mathbb{Z}_2 . This is accomplished using the so called *check matrix representation* or *binary representation* of the stabilizer group. We refer to the appendix as well as to the original work [5] and the treatment in [1] for more details.

5 Preparing states and correcting errors

In the previous section, we have developed the stabilizer formalism to describe quantum error correction codes. Let us now take a more detailed look at how we can actually prepare the logical states and correct errors in that formalism.

Let us first explain how an error correction proceeds when we are using a stabilizer code. Suppose as usual that our stabilizer S is spanned by the independent elements S_1, \dots, S_k , and suppose that $E \in \mathcal{G}_n$ is an error operator.

The first thing to note is that if $|c\rangle \in T_s$ is a code word, and S_i commutes with E , then

$$S_i E|c\rangle = E S_i|c\rangle = E|c\rangle$$

Similarly, if S_i anti-commutes with E , then

$$S_i E|c\rangle = -E S_i|c\rangle = -E|c\rangle$$

Thus the state $E|c\rangle$ is, in any case, in an eigenspace of S_i . Consequently, we can measure all the S_i without changing the state! For each index i , the measurement either gives us $+1$ - this will be the case if and only if E and S_i commute - or -1 - this will be the case if E and S_i anti-commute. Combining the results of the measurements, we obtain a vector e called the *error syndrome*.

As the syndrome only depends on the commutation relation between E and the S_i (and not on the code word), the syndrome gives us a mapping from the set of correctable errors to the set of all possible error syndromes. If this mapping were one-to-one, the error syndrome would tell us exactly which error occurred, and we could apply the inverse of the error operator to correct the state.

However, error correction also works if this mapping is not one-to-one. To see this, suppose we are given two error operators E_1 and E_2 that have the same error syndrome. Thus E_1 commutes with S_i if and only if E_2 does. But this implies that the operator

$$E = E_1^+ E_2$$

commutes with all the S_i . In fact, in the expression

$$E S_i E^+ = E_1^+ E_2 S_i E_2^+ E_1^+$$

moving S_i to the right gives either two times minus one or no sign at all, thus the result will always be S_i . Thus we find that E is in the normalizer $N(S)$. However, by the error correction condition in Theorem 4.1, a product of two errors is either in S or in the complement of $N(S)$. Therefore we find that $E = E_1^+ E_2 \in S$. Thus if one of the two errors occurs for a code word $|c\rangle$, say E_1 , and we apply E_2^+ to correct it, the resulting state will be

$$E_2^+ E_1|c\rangle = E|c\rangle = |c\rangle$$

as if we had applied E_1^+ . The short summary is that even though the error syndrome does not tell us whether the error is E_1 or E_2 , it does not matter - applying any of the inverses will still correct the error.

So once we have an encoded state, we can measure the error syndrome, i.e. measure the S_i (which are Pauli operators and can therefore be measured easily) to obtain all the information that we need to correct the error. But how do we actually prepare an encoded state in the first place?

Again, this turns out to be comparatively easy. We start with some initial state, say the state $|0 \dots 0\rangle$ of the computational basis. Then we measure each of the operators S_i in turn. By the rules of quantum mechanics, the measurement will project the state into an eigenspace of S_i .

If all the eigenvalues are $+1$, we are done - our resulting state will then be in the code space T_S . If one of the eigenvalues is -1 , we need to apply a correction.

For that purpose, we treat the state as if an error with the respective error syndrome had occurred. Thus, we apply an error operator that anti-commutes with those S_i for which the eigenvalue is one and commutes with all the others. This will force the state into a simultaneous $+1$ -eigenstate of all the S_i and therefore into the state T_S (that such an error operator exists can again be shown using the check matrix formalism and linear algebra, see Lemma 1 in [3]).

6 Fault-tolerant quantum computing

Let us briefly recall what we have achieved so far. We have considered an error model where errors occur during a phase which can be thought of as a "transmission" or simply a certain period of time during which a quantum state is stored in a quantum register. We have seen codes that are able to protect against certain types of errors, most notable against single qubit errors. However, our treatment still has a few serious limitations.

First, we have not been able to protect against all errors. Instead, we have only been able to detect and correct single qubit errors. This can be generalized further, but in general we will only be able to protect against errors that affect only a small number of qubits. When we assume that errors on individual qubits occur independently, this is not yet a full protection, but as errors on more than one qubit are less likely in such a model we have been able to reduce the probability of an error.

Further, we have restricted ourselves to errors that occur while we store a state. Obviously, when we act on a state, i.e. when we apply quantum gates, errors can occur as well. Even worse, these errors can propagate. Assume for instance that we have a bit flip in the control qubit of a CNOT gate. If the bit flips from $|0\rangle$ to $|1\rangle$, also the target qubit will flip, thus the error has been duplicated.

Fault-tolerant quantum computing refers to a set of techniques and theoretical results that address these shortcomings. Following [17], the required capabilities can be described as follows.

1. We need to be able to prepare our initial state, i.e. in most cases the logical state $|0\rangle_L$, in a fault-tolerant manner
2. We need to find a way to apply quantum gates to encoded states. The obvious approach - decoding, applying a quantum gate directly and encoding again - could lead to uncontrollable error propagation, as an error in the quantum gate affects all qubits after encoding. Therefore, one typically applies *logical gates* to the encoded states, and we need to be able to do this in a fault-tolerant way
3. After applying a logical gates, we need to be able to detect and correct errors, and again this needs to be done in a fault-tolerant way, as the circuits used for error correction might themselves create errors
4. Finally, our measurements must not introduce an uncontrollable number of new errors and must reach a certain precision

Fault-tolerant quantum computation is a highly non-trivial subject, and we can only provide a few examples to illustrate the types of techniques and arguments used. We refer the reader to the many excellent material on the subject (including original work like [2] or [5] and reviews like [17], [13] and the respective sections in [16] and [1]).

As a first example, let us see how we might detect errors in fault-tolerant way. We again use the three-bit code as an example to keep things simple, even though this code does not protect against phase flip errors.

Let us take a look at the circuit displayed in figure 2 again, that we have used before to measure the error syndrome for the three-bit code. This circuit works, but it is not fault-tolerant. The reason is that the CNOT gate can propagate errors in an uncontrolled way.

To see this, suppose in general that we are given a unitary transformation U and an error E . If the error occurs before applying the gate, the resulting state will be

$$UE|\psi\rangle = UEU^+U|\psi\rangle = (UEU^+)U|\psi\rangle$$

In other words, the result of the operation is the same as if we had applied U to the correct state and then applied the error UEU^+ after passing the gate. Now let us assume that U is a CNOT operation and we apply a phase error to the target bit. Then the resulting error operator will be

$$UZ_2U^+ = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & -1 \end{pmatrix} \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix} = Z_1Z_2$$

In other words, a phase flip error in the target qubit propagates through the CNOT gate as a phase flip error in both involved qubits, target and control.

Returning to our syndrome measurement circuit, we see that an error in the ancillary qubits can propagate into the code qubits. Even worse, an error in one of the ancillary qubits can propagate into *two* of the code qubits, destroying the encoded state beyond repair.

A similar problem occurs for the Shor nine qubit code, and in [2], Shor proposed a solution for the problem (which we replicate here for the three qubit code to keep the diagrams simple, even though the basic idea is the same). Shor replaces each ancillary qubit with two ancillary qubits, each of which is connected to one qubit of the code word by a CNOT. Each pair of ancillas is then prepared in the state

$$\frac{1}{\sqrt{2}}(|00\rangle + |11\rangle)$$

i.e. a superposition of all even two qubit states. This state is kept stable by two or more generally an even number of inversions. However, if we apply an inversion to only one of the qubits, we get a superposition of states with odd parity. Thus we can use the circuit in 4 to determine the error syndrome as a measurement of the ancillary qubits will tell us whether a bit flip error has occurred.

Now this circuit is more tolerant, since an error in one of the ancillary qubits will affect only one of the code word qubits. Thus if that particular case of error can be corrected if only one qubit is affected, we are safe (again note that this is not the case for our particular example, but for Shor's code).




Figure 4: Fault tolerant syndrome measurement

There is of course more that can go wrong. First, the state of the ancillary qubits itself can be faulty. To protect against this, we need a mechanism to validate the ancillary qubits once they have been prepared and repeat the preparation if needed. Second, the measurement of the ancillas can give a wrong result, but we can increase the precision of the measurement by repeating it several times, as discussed in [2]. Overall, these ideas will give us a sufficiently fault tolerant way to measure the error syndrome. Of course the error correction step itself can still go wrong, but the entire procedure will only fail if two independent errors occur (one initial error and one error during the correction) and we will be able to fully recover if only one error occurs.

Having discussed error correction, let us now take a brief look at how we can compute with encoded states, i.e. how we can apply logical gates to the code words. Again, the main challenge will be to do this in a fault-tolerant way, and we will discuss only one example to illustrate the basic idea.

Let us assume we wanted to implement a CNOT gate that operates on two logical qubits using the three qubit code, i.e. on six physical qubits. A straightforward implementation of such a gate could be as follows.

It is rather obvious that this circuit works. If the control state is

$$|1\rangle_L = |111\rangle$$

then each of the CNOT gates will flip a qubit in the target word, and the target state will be flipped logically. If the control state is $|0\rangle_L$, no bit flip occurs and the target state is unchanged.

Unfortunately, this circuit is again not fault tolerant - a bit flip error in the first qubit of the control word will propagate into all three qubits of the target




Figure 5: Not fault tolerant encoded CNOT

word and create an unrecoverable error. The reason for this behaviour is that one qubit in the first code block is connected to more than one - in this case three - qubits in the second code word.

For a better approach, let us take a look at the circuit shown in figure 6.




Figure 6: Fault tolerant encoded CNOT

This circuit also realizes a CNOT on the level of logical states. However, for this circuit, an error in one of the qubits of the first block can only propagate into one qubit in the second block. So if we are able to fully recover from one qubit errors, this circuit works - again as long as only one error occurs. This connectivity property of a circuit is sometimes called *transversality*.

Transversal gates are fault-tolerant, as an error in one qubit can only propagate into one qubit of any other code block, a situation from which we can recover. For many codes, many basic gates can be implemented qubit-wise and

therefore transversally on the level of logical qubits. Unfortunately, a full set of universal quantum gates cannot be constructed in such a way - typically one gate remains for which no transversal implementation can be found. Fortunately, there are other (considerable more complicated) methods available to construct fault-tolerant logical gates. In [2], Shor has given the first fully universal set of fault-tolerant logical gates for his nine qubit code.

7 The threshold theorem

We have now discussed fault-tolerant error correction and fault-tolerant computation on encoded states. Let us now see how this fits into the overall picture of fault-tolerant quantum computation.

Suppose we are given a circuit C that describes a quantum computation. We can think of the overall computation as being organized in individual *time steps*. At each time step, every qubit in the circuit takes part in a basic set of possible *operations* (or, to mention a different term that is sometimes used, *locations*). This can be a gate, the preparation of a qubit in a initial state, a measurement or a "wait step", i.e. the identity operation leaving the involved qubits unchanged. At each time step, a qubit is involved in exactly one of these operations.

Given such a circuit, we can now try to build a fault tolerant circuit performing the same computation as follows. We replace each qubit in the circuit C by a block of qubit encoding a logical qubit. Similarly, we replace each operation in C by the corresponding fault tolerant equivalent, i.e. we replace each gate by the corresponding logical gate, a measurement by a fault tolerant measurement, and a preparation by a fault tolerant state preparation. In addition, we place a (fault tolerant) error correction circuit in each block which is the output of a logical gate. The result is called a *simulation* of the original circuit.

An example is displayed in figure 7, using the graphical language described in [3].

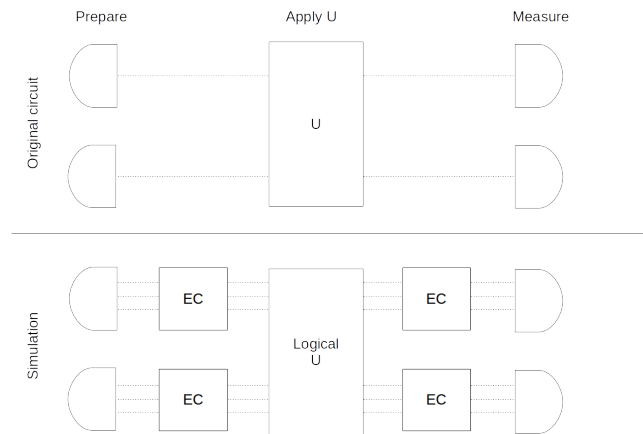


Figure 7: Simulation

The upper part of this diagram shows an original circuit, in which we prepare two qubits in a known state, then apply the unitary transformation U and finally measure both qubits. The result of the replacement procedure described above is displayed in the lower part of the picture, using again the (unrealistic) example of a three qubit code. Here we have replaced each qubit by a corresponding block of three qubits holding the code word. The operation U has been replaced by its logical implementation, acting on the encoded states, and both after the preparation and after applying U , we add a circuit performing error correction.

We have already seen that the circuits used for error correction, state preparation, logical gates and measurement need to be designed with care and in a fault-tolerant way, which essentially means that errors do not propagate in an uncontrolled manner. But even if we do this, it is not clear that the error rate in the simulation is really smaller. The reason is that even though we have added error correction, we have also added many additional operations acting on the physical qubits, i.e. more things that can go wrong. Thus passing to the simulation will introduce additional errors, and it is not a priori clear - and in fact not true - that the error correction will be able to compensate for this so that the net effect is actually an increased reliability.

However, it turns out that you win the race between additional errors that error correction circuitry introduce and the effect of the correction itself provided that the fidelity of the underlying physical operation is below a certain threshold. This is the content of the famous *threshold theorem* ([18], [19], [20]).

The proof of this theorem is long and complicated, so we restrict ourselves to an oversimplified example. Let us again look at the fault-tolerant implementation of a CNOT gate shown in figure 6. To analyze the probability that this circuit fails, let us assume that the only error that can occur is a bit flip error in the target qubit of each physical CNOT gate and that this error occurs with probability p . Thus we assume that error correction, state preparation and measurement work in an idealized, error free way.

Let us now look at the different combinations of faults that can occur during a specific run of the simulation circuit, i.e. a specific *fault path*). A first fault path is given a single bit flip error in one of the target qubits after executing the CNOT operation. However, this affects only one qubit and can therefore be corrected by the error correction circuit. Thus, for this fault path, the simulation works correctly.

Let us now assume a different fault path, given by a fault in target qubit one and target qubit two, whereas the CNOT acting on target qubit three operates correctly. Assuming that the individual faults are statistically independent, the probability for this fault path is

$$p \cdot p \cdot (1 - p) = p^2 - p^3$$

This would introduce an error from which we could not recover, as our code is only able to correct errors in one qubit. But this is not yet the overall probability of failure for our new circuit, as there is a second fault path that we need to consider - we could as well have a failure in the first and the third instead of the first and second qubit. Obviously, the probability for this fault path is the same. And there is a third fault path with two errors, given by an error in the second and third target qubit. These three different fault paths, all with two

errors, correspond to the

$$\binom{3}{2} = 3$$

possibilities to choose two error locations out of the three target qubits. And finally, it might happen that faults occur in all three location. Thus the overall probability to have at least two faults is

$$3(p^2 - p^3) + p^3 = 3p^2 - 2p^3 \leq 3p^2$$

Thus the overall probability that the circuit does not operate correctly does not reduce from p - for the physical CNOT operation without error protection - to p^2 , but to $3p^2$, where the factor 3 reflects the growth of the circuit when passing from the physical circuit to the simulation, i.e. the impact of the additional gates.

Of course this is a very naive model, but it turns out that a similar result can be obtained with more realistic error models - this is the work done in [18] and [19] and more recently in [20] using a different formalism that is better suited to cover the case of a code that can correct only one error as well. The upshot of the discussion is that when passing from a physical circuit with error rate p per operation to the simulation and applying the principles of fault-tolerant quantum computation explained in the previous section, the effective error rate is given by

$$p \mapsto cp^{t+1}$$

where t is the number of errors that the used code can correct and A is a constant that takes the growth of the circuit when passing to the simulation into account and depends on the code and the used fault-tolerant implementations of gates, measurements and state preparations.

Of course this is not necessarily better than p , unless $cp^t < 1$, i.e. unless the fault probability p of the physical operations is below a critical threshold given by

$$p < p_{crit} = c^{-\frac{1}{t}}$$

Now the point is that if we are below this threshold, we can improve further by repeating this procedure. Thus we start with a circuit C_0 , pass to the simulation C_1 which is again made up of physical gates and qubits, and then replace each of these qubits and gates by their encoded equivalents to obtain another circuit C_2 which simulates C_1 and so forth. On the level of codes, this means that in each step, we encode each of the qubits that make up the code word further using the same code - a procedure known as *concatenation*.

Let us see how the probability for a fault develops when we do this. For simplicity, let us consider the case $t = 1$. We have already seen that passing to the simulation amounts to replacing p by

$$cp^2 = p_{crit}^{-1}pp = \epsilon p$$

where $\epsilon = pp_{crit}^{-1} = cp < 1$. If we now iterate this procedure, it is easy to show by induction that after the k -th step, the overall fault probability will be

$$\epsilon^{2^{k+1}-1}p$$

As $\epsilon < 1$, this quickly goes to zero, and we can therefore reach any desired precision, whereas if $\epsilon > 1$, the probability blows up exponentially when we iterate

through this procedure. Thus we have shown that *given physical operations with fault probability less than p_{crit} , we can implement each quantum circuit in a fault-tolerant way with arbitrary precision.*

Unfortunately, the exact value of the threshold is not known. The initial estimate provided in [18] was in the range of $p_{crit} = 10^{-6}$. Later, this was improved by several authors, and is mostly assumed to be in the range between 10^{-4} and 10^{-2} (see [13] for an overview of some results in this direction and [6] for the error threshold for surface codes).

As c is in the order of p_{crit}^{-1} , this implies that we need as many as a few hundred to a few thousand physical gates and qubits to simulate one logical qubit. Note that in general, the overhead depends on the desired accuracy and the level of simulations needed as well as on the factors entering the constant c , i.e. the encoding, the architecture of the logical gates, as well as on the quotient ϵ that measures how close the actual physical error rate is compared to the threshold. Improving the threshold and thereby reducing the overhead remains one of the most active areas of current research in quantum computing and is essential for paving the road towards a large-scale fully fault tolerant universal quantum computer.

A The Pauli matrices

In this section, we fix some notation and recall the most important properties of the Pauli matrices, for which different conventions are in place. We call the matrix

$$X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$$

the bit-flip operator or inversion operator, as it switches the two basis vectors in the standard basis. Similarly, the matrix

$$Z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$$

flips the relative phase of the vectors in the standard basis and is therefore called the phase flip operator. The product

$$Y = ZX = \begin{pmatrix} 0 & 1 \\ -1 & 0 \end{pmatrix}$$

is denoted by Y (whereas other authors use the letter Y differently). The matrices X and Z are hermitian, whereas Y is anti-hermitian. To make contact with the notation usually used in quantum physics, we now introduce the notation

$$\sigma_x = X = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}, \quad \sigma_y = -iY = \begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}, \quad \sigma_z = Z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$$

These matrices are usually called the *Pauli matrices*. They are hermitian and unitary and have square 1. Looking at the linear combinations $\sigma_z \pm 1$ and $\sigma_x \pm i\sigma_y$, it is also not difficult to see that the three Pauli matrices together with

the identity matrix form a complex basis for the set of complex 2x2 matrices, a fact that is exploited heavily in quantum error correction.

Let us now study the multiplicative subgroup generated by the Pauli matrices and the identity. We have the relations

$$\sigma_x \sigma_y = i \sigma_z \quad \text{plus cyclic}$$

and any two Pauli matrices anti-commute. Thus in order to form a group, we will have to add the scalars $\{\pm 1, \pm i\}$. The resulting group that consists of the sixteen matrices

$$\mathcal{G} = \{\pm 1, \pm \sigma_x, \pm \sigma_y, \pm \sigma_z, \pm i, \pm i \sigma_x, \pm i \sigma_y, \pm i \sigma_z\}$$

is called the *Pauli group*. This group is sometimes denoted by \mathcal{G}_1 to indicate that it consists of operators acting on one qubit. For n qubits, we can as well form the Pauli group \mathcal{G}_n that is generated by tensor products of the Pauli matrices acting on one qubit only. In general, the order of the Pauli group is 2^{2n+2} , where the factor $2^{2n} = 4^n$ is given by the pure products of Pauli matrices and the identity and the factor of $2^2 = 4$ is added by $\{\pm 1, \pm i\}$.

Let us now collect a few remarks on stabilizer subgroups of the Pauli group. The first observation is as follows.

LEMMA A.1. *Let S be a subgroup of \mathcal{G} and T be a non-trivial subspace fixed by all elements in S . Then $-1 \notin S$ and S is abelian.*

Proof. It is clear that $-1 \notin S$, as -1 does not stabilize any non-trivial subspace. Now assume that S is not abelian, i.e. that there are elements A and B that do not commute. Any two elements of the Pauli group either commute or anti-commute, thus $AB = -BA$. But if x is some non-zero element of T , this implies

$$x = (AB)(x) = -(BA)(x) = -x$$

which is a contradiction. \square

So in order to stabilize a non-trivial subspace, the subgroup S needs to be abelian and must not include -1 . It is instructive to note that a subgroup of \mathcal{G}_n that does not contain -1 is automatically abelian. In fact, given two elements a, b of such a group, both elements square to one (as every element of the Pauli group squares to one or minus one and minus one is not allowed). Suppose that the elements anti-commute (again, any two elements of the Pauli group either commute or anti-commute). Then

$$ab = -ba$$

which, by multiplying from the right first by a and then by b , implies

$$abab = -1$$

which is a contradiction as minus one is not in our subgroup.

Let us study the structure of such subgroups of \mathcal{G}_n in more detail. First, every element of \mathcal{G}_n has square ± 1 . As we assume that -1 is not in S , each element of S therefore squares to one, i.e. each element of S has order two. Further, S is a finite abelian group. By the structure theory for finite abelian groups, we can therefore conclude that

$$S = \mathbb{Z}_2 \oplus \mathbb{Z}_2 \cdots \oplus \mathbb{Z}_2$$

Let us denote the number of factors by l . Then the order of S is 2^l . As t elements of order t can at most generate a subgroup of size 2^t , it also follows that the number of factors l is equal to the minimum number of generators that we need to generate the group S . Thus any set of independent generators will have length l .

Under these conditions, one can also relate the dimension of the subspace to the number of generators of the group. In fact, we have the following

LEMMA A.2 (see [4] or [16], section 11.4.2). *Suppose that S is an abelian subgroup of the Pauli group \mathcal{G}_n such that S is generated by k independent elements and $-1 \notin S$. Then the subspace T_S of all vectors that are fixed by all elements of S has dimension 2^{n-k} . The projection onto this subspace is given by the product*

$$\frac{1}{2^k} \prod_{i=1}^k (1 + S_i)$$

Proof. Let us first prove the following claim: if S_1, \dots, S_k are independent generators of S , the projection onto T_S is given by the product

$$\frac{1}{2^k} (1 + S_k)(1 + S_{k-1}) \cdots (1 + S_1)$$

To see this, we use induction on k . For any $r \leq k$, let us denote by T_r the subspace which is fixed by all generators S_1, \dots, S_r . As S_{r+1} commutes with every S_i , it is clear that S_{r+1} is mapping T_r to itself. Further, we have

$$T_{r+1} = \{|\psi\rangle \in T_r \text{ such that } S_{r+1}|\psi\rangle = |\psi\rangle\}$$

In other words, T_{r+1} is the $+1$ eigenspace of S_{r+1} restricted to T_r .

Let P_r be the orthogonal projection onto T_r . To prove our claim by induction, we have to show that the operator

$$\frac{1}{2}(1 + S_{r+1})P_r$$

is the orthogonal projection onto T_{r+1} , i.e. the unique idempotent hermitian operator with range T_{r+1} . Now each S_{r+1} is hermitian (as it is unitary with square one), and P_r is hermitian, so this operator is clearly hermitian. Using the induction hypothesis, it is also clear that P_r commutes with S_{r+1} and a short calculation using this shows that our operator squares to one. So we only have to show that its range is equal to T_{r+1} .

To do this, let us first show that its range is contained in T_{r+1} . As S_{r+1} acts on T_r , a vector of the form

$$|x\rangle = \frac{1}{2}(1 + S_{r+1})P_r|y\rangle$$

is of course contained in T_r . Now

$$S_{r+1}(1 + S_{r+1}) = S_{r+1} + 1 = 1 + S_{r+1}$$

and hence

$$S_{r+1}|x\rangle = S_{r+1}\frac{1}{2}(1 + S_{r+1})P_r|y\rangle = \frac{1}{2}(1 + S_{r+1})P_r|y\rangle = S_{r+1}|x\rangle$$

which shows that the range is contained in T_r . To show the converse, assume we are given an element in $|x\rangle \in T_{r+1}$. Then S_{r+1} fixes this element, and the same is true for P_r , so we have

$$|x\rangle = \frac{1}{2}(1 + S_{r+1})P_r|x\rangle$$

showing that $|x\rangle$ is in the range. This completes the proof of our claimed formula for the projection P_{r+1} .

Now let us turn to proving the main statement of the Lemma. Again, we do this by induction. Suppose we have already shown that $\dim T_r = 2^{n-r}$. We then only have to show that the eigenspace of S_{r+1} with eigenvalue one cuts out half of T_r to complete the induction step. This is equivalent to showing that the dimensions of the plus one and minus one eigenspaces of $S_{r+1}P_r$ are equal, which again is (as one can see in a basis of eigenvectors) equivalent to the statement that the trace of $S_{r+1}P_r$ is zero. But by the formula for the projection, we see that

$$S_{r+1}P_r = S_{r+1}(1 + S_r)(1 + S_{r-1})\cdots(1 + S_1)$$

This is a sum of terms each of which is a product of elements of \mathcal{G}_n different from the identity. Each such product has trace zero, and therefore the trace of the sum is zero and our proof is complete. \square

LEMMA A.3. *Assume that S is an abelian subgroup of the Pauli group \mathcal{G}_n and that the subspace*

$$T_S = \{|\psi\rangle \text{ such that } s|\psi\rangle = |\psi\rangle \forall s \in S\}$$

is non-trivial. Then the subgroup S is equal to the full stabilizer of T_S , i.e. if g is an element of \mathcal{G} that fixes all elements of T_S , then $g \in S$.

Proof. Pick a set of l elements s_1, \dots, s_l that are independent and generate S . Let $\bar{S} = S(T)$ denote the full stabilizer of T_S , i.e.

$$\bar{S} = \{g \in \mathcal{G}_n \text{ such that } gx = x \forall x \in T_S\}$$

Clearly, $S \subset \bar{S}$. We want to show that $\bar{S} = S$. Suppose that this is not the case, i.e. that we can find an element $s \in \bar{S}$ which is not already in S . Let us then consider the group S' generated by the set s, s_1, \dots, s_l . Clearly, $T_S = T_{S'}$, and the set of generators s, s_1, \dots, s_l is independent. But given Lemma A.2, this would imply that

$$\dim T_S = 2^{n-l} = \dim T_{S'} = 2^{n-l+1}$$

which is a contradiction. \square

B The check matrix formalism

In this section, we will shortly describe the check matrix formalism and show for a few examples how it can be applied to stabilizer codes. Again, our setup will be as follows. We are given n physical qubits and $k < n$ independent elements S_1, \dots, S_k of the n -dimensional Pauli group \mathcal{G}_n which commute. We consider the abelian subgroup $S = \langle S_1, \dots, S_k \rangle$ generated by these elements and assume that $-1 \notin S$. By what we have seen in the previous section, we therefore know that the code space T_S is non-trivial and has dimension 2^{n-k} .

Now assume that we are given an element g of the Pauli group \mathcal{G}_n . As, up to a phase, σ_y is a product of σ_x and σ_z , we can write this as

$$g = \mu(\sigma_x^1)^{a_1}(\sigma_z^1)^{b_1} \dots (\sigma_x^n)^{a_n}(\sigma_z^n)^{b_n}$$

with an overall phase factor $\mu \in \{\pm 1, \pm i\}$. To such an element, we then associate the vector (following the conventions in [5] which are similar to ours)

$$r(g) = (a, b) \in \mathbb{Z}_2^n \times \mathbb{Z}_2^n = \mathbb{Z}_2^{2n}$$

Thus the vector a has a one at position i if σ_x occurs in the i -th qubit, and b has a one at position i if σ_z occurs at the i -th qubit. The vector a is sometimes called the x -vector and b is sometimes called the z -vector of g , together they form the $2n$ -dimensional binary vector $r(g)$.

The mapping from \mathcal{G}_n to \mathbb{Z}_2^{2n} is easily seen to be a group homomorphism, i.e. $r(gg') = r(g) + r(g')$. Furthermore, the kernel is exactly the subgroup of order four generated by i . In particular the restriction of r to S is one-to-one, as -1 is not in S and therefore i and $-i$ are not in S as well.

This homomorphism is extremely useful to relate statements about elements of the Pauli group to linear algebra over the field \mathbb{Z}_2 . As an example, let us prove the following fact.

LEMMA B.1. *Suppose S_1, \dots, S_k are elements of the Pauli group \mathcal{G}_n such that $-1 \notin \langle S_1, \dots, S_k \rangle$. Then the S_i are independent if and only if the binary vectors $r(S_i)$ are linearly independent.*

Proof. First, suppose that the $r(S_i)$ are linearly independent and assume that we have a relation among the S_i , which (as the S_i square to one and commute) we can always arrange to be in the form

$$1 = \prod_i S_i^{\alpha_i}$$

with $\alpha_i \in \mathbb{Z}_2$. Applying the r -homomorphism yields

$$0 = r(1) = \sum_i \alpha_i r(S_i)$$

As the $r(S_i)$ were assumed to be linearly independent, this implies that all coefficients α_i are zero.

Conversely, suppose we are given a relation between the $r(S_i)$. Then the same calculation shows that

$$0 = r\left(\prod_i S_i^{\alpha_i}\right)$$

which shows that this product is in the kernel of r , i.e. it shows that

$$\prod_i S_i^{\alpha_i} \in \{\pm 1, \pm i\}$$

But as -1 is not in the group generated by the S_i , the only value that this product can take is one. Therefore this gives us a relation between the S_i , and as those are independent, we can again conclude that the α_i all vanish. \square

Now let us try to find a condition on the $r(g)$ that tells us whether two elements commute or anti-commute. Suppose we are given two group elements g and g' and

$$\begin{aligned} r(g) &= (a, b) \\ r(g') &= (c, d) \end{aligned}$$

Then the binary scalar product ad measures in how many qubits a σ_x in g and a σ_z in g' pair up. Each of these pairs gives us a factor minus one if we commute g' by g . Similarly, cb measures in how many qubits a σ_x in g' meets a σ_z in g . Thus we see that g and g' commute if and only if

$$ad + bc = 0$$

and anti-commute if and only if

$$ad + bc = 1$$

This formula defines a symplectic scalar product on the vector space \mathbb{Z}_2^{2n} , which is of course simply the standard symplectic product defined on any even-dimensional vector space. To write down this product, we can introduce the $2n \times 2n$ matrix

$$\Lambda = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}$$

which switches a and b when acting on a vector (a, b) . Thus we have shown:

LEMMA B.2. *Two elements g and g' of the Pauli group commute if and only if*

$$r(g')\Lambda r(g)^T = 0$$

Finally, it is sometimes useful to understand how an operator g acts on an element $|a\rangle$ of the standard basis in the space of n qubits. To see this, write a as a binary string

$$|a\rangle = |a_1 a_2 \dots a_n\rangle$$

Now suppose we are given an element g of the Pauli group, again written without any σ_y terms. Then each σ_x acts by flipping a qubit, whereas each σ_z acts by adding a phase. Thus if the i -th position of $r(g)$ is one, a bit is flipped, otherwise a phase is picked up. Let us denote by $r_1(g)$ the projection of $r(g)$ to the first n coordinates. Then, in a hopefully readable notation,

$$g|a\rangle = \mu|a \oplus r_1(g)\rangle$$

with a phase μ , where the \oplus denotes bitwise addition.

After these preliminaries, we can now introduce check matrices. Suppose we are in the standard situation considered in this section, i.e. we have an abelian group S with independent generators S_1, \dots, S_k that stabilizes a non-trivial subspace T_S .

For each of the generators S_i , we can again consider the $2n$ -dimensional binary vector $r(S_i)$. Let us combine these vectors into a matrix \mathcal{C} called the *check matrix*. Thus if

$$r(S_i) = (a_i, b_i)$$

then the check matrix is the matrix

$$\mathcal{C} = (A \mid B)$$

with k rows and $2n$ columns, where we use the vertical bar in the middle to separate the first n columns from the second n columns. As the elements S_i are independent, we know from our previous discussion that the matrix has full rank r .

To illustrate how the check matrix can be used to translate group theoretical questions into ordinary linear algebra, let us prove the following statement.

LEMMA B.3. *The normalizer $N(S)$ has $4 \cdot 2^{2n-k}$ elements.*

Proof. We know that an element g is in the normalizer of S if and only if it commutes with every generator S_i , i.e. if

$$r(S_i)\Lambda r(g)^T = 0$$

for all i . Now the $r(S_i)$ are the rows of \mathcal{C} . Combining these k equations is therefore equivalent to the condition

$$\mathcal{C}\Lambda r(g)^T = 0$$

As Λ is an isomorphism, the kernel of $\mathcal{C}\Lambda$ has the same dimension as the kernel of \mathcal{C} . But the dimension of this kernel is $2n - k$, as the rank of the check matrix is k .

So the normalizer is the pre-image of a linear subspace of dimension $2n - k$ under the homomorphism r . The kernel of this homomorphism is the group $Z = \{\pm 1, \pm i\}$ which is in the center of the normalizer. This gives an additional factor four and shows that the cardinality of the normalizer is $4 \cdot 2^{2n-k}$ as claimed. \square

But we can do even more when we bring the check matrix into a standard form. In fact, we can start to apply the basic steps of a Gauss-Jordan elimination to it, more specifically to its left side! Let us see how these steps are reflected in the group S .

First, we can of course re-label the qubits at will. This will simply exchange two columns of the left side (and the corresponding columns on the right side as well). Next, we can re-label the generators. This simply corresponds to a permutation of the rows. And finally, we can pick an index i and multiply all generators except S_i by S_i , i.e. pass to the set $\{S_1 S_i, S_2 S_i, \dots, S_i, \dots, S_k S_i\}$. This corresponds to adding row i to all other rows of the matrix. Thus all operations in the Gauss-Jordan algorithm (over \mathbb{Z}_2) are reflected by operations on the group level that do not change the stabilizer group S .

Thus, starting with any check matrix, we can apply Gauss-Jordan elimination steps, i.e. change the set of generators by the operations explained above, to the left hand side of the matrix and end up with a check matrix in the following structure.

$$\mathcal{C} = \left(\begin{array}{cc|cc} 1 & A & B & C \\ 0 & 0 & D & E \end{array} \right) \quad (2)$$

Here the unit matrix in the upper left corner has dimension r , where r is the rank of the left hand side of the check matrix.

As an example how useful this representation can be, let us now see how we can systematically construct vectors in the code space T_S , i.e. vectors invariant under S . The obvious approach which readers familiar with group theory will immediately recognize is averaging, i.e. starting with some element, say $|0 \dots 0\rangle$, and then looking at the sum

$$|x\rangle = \sum_{s \in S} s|0 \dots 0\rangle$$

As multiplying by an element in S simply reorders the sum, it is clear that this element is invariant. However, it is not clear that it is not zero (and we have not used the necessary condition that $-1 \notin S$ so far). However, with the standard assumptions that we usually make, this is true, and the check matrix standard form can be used to prove it.

LEMMA B.4. *Suppose that the abelian group S has a non-trivial invariant subspace T_S . Then the invariant element*

$$\sum_{s \in S} s|0 \dots 0\rangle$$

is not zero.

Proof. So far we have not chosen any generators. Let us do this now, i.e. we pick independent generators S_1, \dots, S_k . As we assume that the code space is non-trivial, we already know that $k < n$ and that $-1 \notin S$. By the arguments above, we can also assume that the generators are chosen such that the check matrix is in the standard form (2) described above.

We already know (and there the condition that $-1 \notin S$ comes into play) that the group S is a product $(\mathbb{Z}_2)^k$. Using this, it is not difficult to show that

$$\sum_{s \in S} s|0 \dots 0\rangle = \frac{1}{2^k} \prod_{i=1}^k (1 + S_i) |0 \dots 0\rangle$$

and therefore be expressed by the generators alone. Now let again r denote the rank of the left hand side of the check matrix. Then all S_i with $i > r$ do not contain any σ_x matrices, as the rows of the check matrix below index r are zero in the first n positions. As σ_z acts trivially on $|0 \dots 0\rangle$, we can ignore these factors (as they only give a factor two) and see that

$$\sum_{s \in S} s|0 \dots 0\rangle = \frac{1}{2^k} \prod_{i=1}^r (1 + S_i) |0 \dots 0\rangle$$

which expresses our element by the first r generators only (these generators are called type I generators in [5]). Reversing the order, we can write the same expression as

$$\sum_{s \in S} s |0 \dots 0\rangle = \frac{1}{2}(1 + S_r) \dots \frac{1}{2}(1 + S_1) |0 \dots 0\rangle$$

Now let us take a look at the check matrix in the standard form to see how S_1 acts. We focus on the first r bits and write a binary vector of length n as

$$|x; ?\rangle$$

with a binary string x of length r and the question mark indicating that we do not care about the other positions. Then, obviously

$$S_1 |0 \dots 0\rangle = \mu_1 |10 \dots 0; ?\rangle$$

as S_1 only flips the first qubit, where μ_1 is a phase factor. Thus

$$(1 + S_1) |0 \dots 0\rangle = (|0 \dots 0; ?\rangle + \mu_1 |10 \dots 0; ?\rangle)$$

If we now apply S_2 to this, then S_2 will only flip the second qubit and let the bits among the first r bits alone (and it might act on the qubits after position r). Thus

$$S_2(1 + S_1) |0 \dots 0\rangle = (\mu_2 |01 \dots 0; ?\rangle + \mu_3 |110 \dots 0; ?\rangle)$$

with new phase factors μ_2, μ_3 . Therefore

$$\begin{aligned} (1 + S_2)(1 + S_1) |0 \dots 0\rangle = & |0 \dots 0; ?\rangle + \mu_1 |10 \dots 0; ?\rangle \\ & + \mu_2 |01 \dots 0; ?\rangle + \mu_3 |110 \dots 0; ?\rangle \end{aligned}$$

which already has four components. Using again the standard shorthand notation (replacing binary strings by their decimal value), we could write this as well as

$$(1 + S_2)(1 + S_1) |0 \dots 0\rangle = |0; ?\rangle + \mu_1 |1; ?\rangle + \mu_2 |2; ?\rangle + \mu_3 |3; ?\rangle$$

Now we can continue like this, exploiting the fact that each new S_i will only act on bits among the first r bits which all the previous operators have left alone. We therefore see that

$$\sum_{s \in S} s |0 \dots 0\rangle = \frac{1}{2^k} \sum_{i=0}^r \mu_i |i; ?\rangle$$

where all the μ_i are non-zero. Thus this is a sum of non-zero coefficients times different vectors of the standard basis and therefore non-trivial. \square

One can use similar arguments and a more refined standard form to obtain much more information on the structure of the normalizer $N(S)$. This is carried out in section 4.1 of [5]. The main result is that it is possible to find a set of $n-k$ group elements \bar{X}_i, \bar{Z}_i in the normalizer $N(S)$ with the following properties.

1. All the \bar{X}_i commute with each other and with any element in S

2. All the \bar{Z}_i commute with each other and with any element in S
3. \bar{X}_i and \bar{Z}_j commute for $i \neq j$
4. \bar{X}_i and \bar{Z}_i anti-commute
5. The \bar{Z}_i operators are built of σ_z matrices only

Clearly, the \bar{Z}_i are hermitian, as they are a product of σ_z -operators. As the properties above are not changed if we multiply any of the \bar{X}_j with i , we can also assume that the \bar{X}_i are hermitian. Further, none of the \bar{X}_i or \bar{Z}_i can be in S , as each of them anti-commutes with some element in $N(S)$ (\bar{Z}_i anti-commutes with \bar{X}_i).

Let us see what else these properties imply. First, let us revisit the (non-zero) state which is obtained by averaging over the full group - we will denote this state by $|\bar{0}\rangle$ in the sequel:

$$|\bar{0}\rangle = \sum_{s \in S} s |0 \dots 0\rangle$$

As any \bar{Z}_i can be commuted past the $s \in S$ and acts trivially on the state $|0 \dots 0\rangle$ (being composed of σ_z 's only), we find that

$$\bar{Z}_i |\bar{0}\rangle = \bar{Z}_i \sum_{s \in S} s |0 \dots \rangle = \sum_{s \in S} s \bar{Z}_i |0 \dots \rangle = |\bar{0}\rangle$$

So $|\bar{0}\rangle$ is a $+1$ eigenstate of all the \bar{Z}_i , very similar to the state $|0\rangle$ being an eigenstate of σ_z . Now we can apply the type of arguments that physically inclined readers will know from the algebraic treatment of the quantum harmonic oscillator to learn more about the action of the \bar{X}_i on the "ground state" $|\bar{0}\rangle$. For instance

$$\langle \bar{0} | \bar{X}_i | \bar{0} \rangle = \langle \bar{0} | \bar{X}_i \bar{Z}_i | \bar{0} \rangle = -\langle \bar{0} | \bar{Z}_i \bar{X}_i | \bar{0} \rangle = -\langle \bar{0} | \bar{X}_i | \bar{0} \rangle$$

which implies that $\bar{X}_i |\bar{0}\rangle$ is orthogonal to $|\bar{0}\rangle$ (and of course not zero, as \bar{X}_i is unitary). Similar arguments show that given any two different bit-strings (a_1, \dots, a_n) and (b_1, \dots, b_n) , the two vectors

$$\bar{X}_1^{a_1} \dots \bar{X}_n^{a_n} |\bar{0}\rangle$$

and

$$\bar{X}_1^{b_1} \dots \bar{X}_n^{b_n} |\bar{0}\rangle$$

are orthogonal. Thus the encoding

$$(a_1, \dots, a_{n-k}) \mapsto \bar{X}_1^{a_1} \dots \bar{X}_n^{a_n} |\bar{0}\rangle$$

creates an orthonormal basis of the vector space T_S , presenting this vector space as a tensor product of $n-k$ copies of a two-dimensional Hilbert space and thereby identifying the quotient $N(S)/S$ with the Pauli group \mathcal{G}_{n-k} .

Thus the action of $N(S)/S$ on T_S can be identified with the action of the Pauli group on a standard quantum register with $n-k$ qubits. We have therefore not only encoded $n-k$ logical qubits in n physical qubits, but also identified unitary operators, i.e. quantum gates, that act on the encoded qubits similar

to the action of the Pauli matrices on the physical qubits. This is one of the advantages of the stabilizer formalism - it does not only provide a code, but also a set of Pauli operators on the encoded states.

To conclude this section, let us now verify that instead of starting with a group S defined by independent generators S_i , we could as well choose a check matrix as starting point and construct the code from there. Specifically, we have the following result.

LEMMA B.5. *Assume that \mathcal{C} is a $r \times 2n$ matrix with independent rows such that the symplectic product of any two rows is zero. Then there are commuting and independent hermitian group elements S_1, \dots, S_k with $S_i^2 = 1$ such that the group S generated by the S_i does not contain -1 .*

Proof. First, let us choose group elements S_i such that $r(S_i)$ is the i -th row of the given check matrix \mathcal{C} . Then, according to our assumptions about this matrix, the S_i commute. Being an element of the Pauli group, each S_i is either hermitian or anti-hermitian, and each S_i is unitary. Thus the condition $S_i^2 = 1$ is equivalent to S_i being hermitian. If any of the S_i is not hermitian, then we can replace it by iS_i , and therefore we can find the S_i such that they are all hermitian and have square $+1$. Let S denote the group generated by the S_i which is clearly abelian.

The only thing that is left to show is that $-1 \notin S$. To see this, consider the map

$$s: \mathbb{Z}_2^k \rightarrow S$$

that maps a vector $a = (a_1, \dots, a_k)$ to the group element

$$s(a) = \prod_i S_i^{a_i}$$

As the S_i commute and have square $+1$, this map is well-defined and a group homomorphism. By the definition of S as the group generated by the S_i , it is onto.

Let us now assume that $-1 \in S$. Then we could find a vector $a \in \mathbb{Z}_2^k$ such that $s(a) = 1$. Applying the homomorphism r on both sides yields

$$0 = r(-1) = r(s(a)) = \sum_i a_i r(S_i)$$

But we have assumed that the $r(S_i)$ are linearly independent, so this implies that $a = 0$. But $s(0) = 1 \neq -1$, so that we have found a contradiction and our proof is complete. \square

References

- [1] M.A. Nielsen, I.L. Chuang, *Quantum Computation and Quantum Information*, Cambridge University Press, Cambridge, 2010
- [2] P. Shor, *Polynomial-Time Algorithms for Prime Factorization and Discrete Logarithms on a Quantum Computer*, SIAM J.Sci.Statist.Comput. Vol. 26 Issue 5 (1997), pp 1484–1509, available as arXiv:quant-ph/9508027v2

- [3] D. Gottesman, *An Introduction to Quantum Error Correction and Fault-Tolerant Quantum Computation*, arXiv:0904.2557v1
- [4] D. Gottesman, *A Class of Quantum Error-Correcting Codes Saturating the Quantum Hamming Bound*, arXiv:quant-ph/9604038
- [5] D. Gottesman, *Stabilizer Codes and Quantum Error Correction*, Caltech Ph.D. Thesis, available as arXiv:quant-ph/9705052
- [6] A.G. Fowler, M. Mariantoni, J.M. Martinis, A.N. Cleland, *Surface codes: Towards practical large-scale quantum computation*, arXiv:1208.0928v2
- [7] K. Fujii, *Quantum Computation with Topological Codes - from qubit to topological fault-tolerance*, arXiv:1504.01444
- [8] D. Dieks, *Communication by EPR devices*, Physics Letters A, Vol. 92 Issue 6 (1982), 271–272
- [9] W.K. Wootters, W.H. Zurek, *A single quantum cannot be cloned*, Nature Vol. 299 (1982), 802–803
- [10] A.M. Steane, *A tutorial on quantum error correction*, Proceedings of the International School of Physics Enrico Fermi, course CLXII, Quantum Computers, Algorithms and Chaos, G. Casati, D. L. Shepelyansky and P. Zoller, eds., pp. 132 (IOS Press, Amsterdam 2006)
- [11] A.M. Steane, *Multiple Particle Interference and Quantum Error Correction*, arXiv:quant-ph/9601029
- [12] P. Shor, *Scheme for reducing decoherence in quantum computer memory*, Phys. Rev. A Vol. 52 (1995) Issue 4, pp. 2493–2496
- [13] S. J. Devitt, W. J. Munro, Kae Nemoto, *Quantum Error Correction for Beginners*, arXiv:0905.2794v4
- [14] E. Knill, R. Laflamme, *A Theory of Quantum Error-Correcting Codes*, arXiv:quant-ph/9604034
- [15] R. Laflamme, C. Miquel, J.P. Paz, W.H. Zurek, *Perfect Quantum Error Correction Code*, arXiv:quant-ph/9602019
- [16] E. Rieffel, W. Polak, *Quantum computing - a gentle introduction*, MIT Press, Cambridge 2011
- [17] J. Kempe, *Approaches to quantum error correction*, Séminaire Poincaré 1 (2005), pp. 65–93, available as arXiv:quant-ph/0612185
- [18] D. Arahonov, M. Ben-Or, *Fault Tolerant Quantum Computation with Constant Error*, arXiv:quant-ph/9611025
- [19] E. Knill, R. Laflamme, W.H. Zurek, *Threshold accuracy for quantum computation*, arXiv:quant-ph/9610011,
- [20] , P. Aliferis, D. Gottesman, J. Preskill, *Quantum accuracy threshold for concatenated distance-3 codes*, arXiv:quant-ph/0504218v3